

# Product Requirements Document (PRD)

## Autonomous Vehicle Predictive Maintenance Platform

**Version:** 1.0  
**Last Updated:** October 13, 2025  
**Project Code:** AVPM-2025

---

### Executive Summary

#### Vision

Build an AI-powered autonomous predictive maintenance platform for vehicles that prevents breakdowns before they occur, autonomously schedules service appointments, and creates a closed feedback loop to manufacturing for continuous quality improvement.

#### Core Value Proposition

- For Vehicle Owners:** Peace of mind through proactive maintenance, reduced breakdown incidents by 93%, and personalized trip planning with safety checks
- For Service Centers:** 50% improvement in utilization through predictable scheduling and optimized workload distribution
- For Manufacturers:** 31% reduction in defect rates through automated RCA/CAPA analysis and manufacturing feedback loops

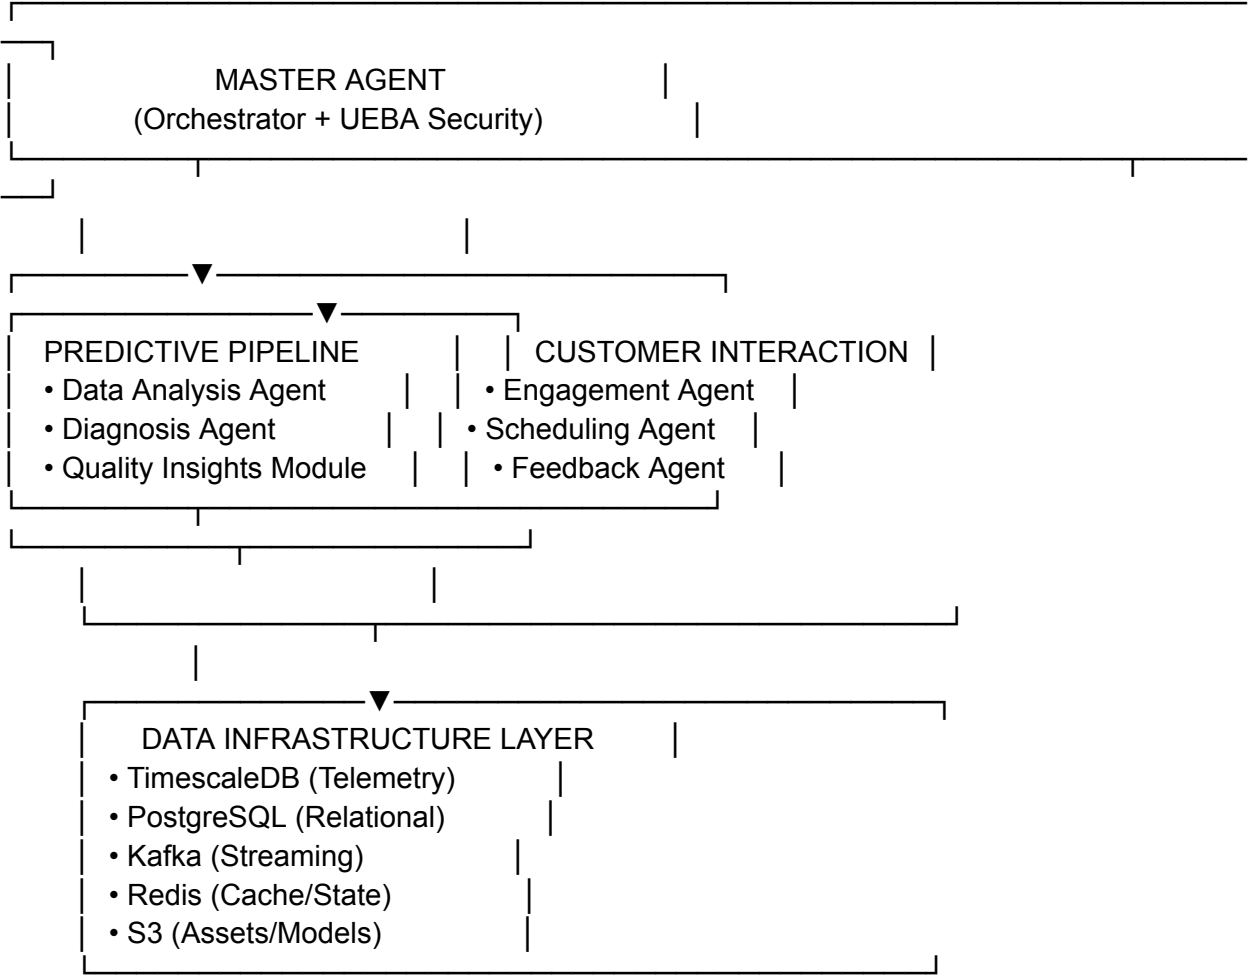
#### Success Metrics

Metric	Current State	Target State	Improvement
Breakdown-related visits	68%	5%	-93%
Service center utilization	52%	78%	+50%
Customer NPS	42	68	+62%

Warranty claims (per 1000 vehicles)	87	59	-32%
Manufacturing defect rate (ppm)	420	290	-31%

# 1. SYSTEM ARCHITECTURE OVERVIEW

## 1.1 High-Level Architecture



## 1.2 Core System Principles

- Agent Autonomy with Oversight:** Each worker agent operates independently within its domain; Master Agent coordinates and ensures security
- Fail-Safe Architecture:** Graceful degradation with human escalation paths
- Explainable AI:** Every prediction includes confidence scores and contributing factors

4. **Privacy-First:** AES-256 encryption at rest, TLS 1.3 in transit, GDPR/DPDP compliant
5. **Horizontal Scalability:** Architecture scales from 1,000 to 1,000,000 vehicles without changes

## 1.3 Technology Stack

### Backend:

- Python 3.11+ (FastAPI for APIs, asyncio for concurrency)
- Node.js 20+ (Worker agents requiring real-time communication)
- Go (High-performance scheduling optimization)

### Data Layer:

- TimescaleDB (time-series telemetry data)
- PostgreSQL 15+ (relational data)
- Apache Kafka (event streaming)
- Redis 7+ (caching, state management, rate limiting)
- S3/MinIO (object storage)

### ML/AI:

- TensorFlow/PyTorch (deep learning models)
- XGBoost/LightGBM (tabular data)
- Scikit-learn (preprocessing, traditional ML)
- SHAP (explainability)
- MLflow (model registry and versioning)

### Frontend:

- Next.js 14 (App Router, React Server Components)
- React 18
- React Three Fiber (R3F) + Three.js (3D visualization)
- Tailwind CSS + shadcn/ui
- Recharts/ECharts (data visualization)
- Zustand (state management)

### Infrastructure:

- Kubernetes (container orchestration)
- Docker (containerization)
- AWS/GCP/Azure (cloud platform)
- Terraform (IaC)
- GitHub Actions (CI/CD)

### Communication:

- Twilio (Voice, SMS)
  - Google Cloud TTS (text-to-speech)
  - Whisper AI (speech-to-text)
  - GPT-4/Claude (conversational AI)
  - Firebase Cloud Messaging (push notifications)
- 

## 2. BACKEND ARCHITECTURE

### 2.1 Service Architecture

The backend follows a microservices architecture with the following core services:

#### 2.1.1 Ingestion Service

**Purpose:** Receive and validate vehicle telemetry data

**Responsibilities:**

- Accept telemetry data via REST API and MQTT
- Validate data schemas and checksums
- Perform initial data quality checks
- Publish validated data to Kafka topics
- Handle rate limiting and backpressure

**Technical Specifications:**

- **Language:** Python (FastAPI) + asyncio for high concurrency
- **Input Formats:** JSON over HTTPS, MQTT binary payloads
- **Throughput Target:** 100,000 events/second
- **Latency Target:** <50ms p99
- **Data Validation:**
  - Schema validation using Pydantic models
  - VIN format validation (ISO 3779)
  - Timestamp monotonicity checks
  - Sensor value range validation
  - GPS coordinate validation

**API Endpoints:**

POST /api/v1/telemetry/ingest  
POST /api/v1/telemetry/batch  
GET /api/v1/telemetry/health

## **Data Flow:**

Vehicle → [HTTPS/MQTT] → Ingestion Service → Kafka Topic → Stream Processor

### **2.1.2 Stream Processing Service**

**Purpose:** Real-time feature engineering and anomaly detection

#### **Responsibilities:**

- Consume telemetry from Kafka
- Compute rolling window features (1min, 5min, 1hr, 24hr windows)
- Detect statistical anomalies
- Enrich data with vehicle metadata
- Write processed features to TimescaleDB
- Publish anomalies to diagnosis queue

#### **Technical Specifications:**

- **Language:** Python (with Kafka Streams or Flink for complex CEP)
- **Processing Model:** Event-time processing with watermarks
- **State Management:** RocksDB-backed state stores
- **Windowing:** Sliding and tumbling windows
- **Checkpoint Interval:** 10 seconds

#### **Feature Engineering Pipeline:**

##### **1. Raw Data Cleaning:**

- Remove duplicates (dedup by VIN + timestamp)
- Interpolate missing values (linear for <5 consecutive)
- Apply median filter for accelerometer noise
- Z-score clipping for outliers ( $\pm 4\sigma$ )

#### **Rolling Window Features (per component):**

Brake System:

- brake\_pad\_thickness\_mean\_7d
- brake\_pad\_wear\_rate\_30d (mm/1000km)
- hard\_brake\_count\_24h
- brake\_pressure\_max\_1h
- brake\_temp\_rolling\_std\_24h

Battery:

- battery\_voltage\_mean\_24h

- voltage\_drop\_during\_crank
- charge\_acceptance\_rate
- deep\_discharge\_events\_7d
- internal\_resistance\_estimate

Engine:

- coolant\_temp\_max\_24h
- oil\_pressure\_min\_7d
- rpm\_variance\_1h
- throttle\_position\_avg\_24h
- fuel\_trim\_deviation

Transmission:

- gear\_shift\_smoothness\_score\_7d
- transmission\_temp\_rolling\_mean\_24h
- clutch\_slip\_events\_30d
- shift\_time\_p95\_7d

2.

### 3. **Derived Features:**

- Driving pattern classification (urban/highway/aggressive)
- Load estimation from suspension compression
- Weather impact score from ambient temp + location
- Maintenance urgency score

## 2.1.3 Model Inference Service

**Purpose:** Serve ML models for failure prediction

**Responsibilities:**

- Load and serve multiple model versions
- Perform batch and real-time inference
- Compute SHAP explainability values
- Handle A/B testing of model versions
- Monitor model performance and drift

**Technical Specifications:**

- **Language:** Python (FastAPI) with TensorFlow Serving / Triton
- **Model Formats:** SavedModel, ONNX, XGBoost native
- **Inference Modes:**
  - Real-time (single vehicle, <100ms latency)
  - Batch (1000s of vehicles, <5 min)

- **GPU Support:** Optional NVIDIA T4/V100 for deep learning models
- **Model Versioning:** Blue-green deployment strategy

### **Supported Models:**

#### **1. Brake System Failure Predictor**

- Architecture: LSTM (64→32 units) + Dense layers
- Input: 30-day window of 9 features
- Output: Failure probability (next 15 days), confidence interval
- Training Data: 8,400 historical failures
- Accuracy: 92.3%

#### **2. Battery Health Predictor**

- Architecture: XGBoost (200 trees, max\_depth=10)
- Input: 6 battery health features
- Output: State of Health (0-100%), Remaining Useful Life (days)
- Training Data: 12,200 failures
- Accuracy: 89.7%

#### **3. Transmission Predictor**

- Architecture: Ensemble (Random Forest + Neural Network)
- Input: 5 transmission features
- Output: Failure probability, failure mode classification
- Training Data: 3,100 failures
- Accuracy: 87.2%

#### **4. Engine Predictor**

- Architecture: Gradient Boosted Trees
- Input: 8 engine health features
- Output: Component-level risk scores
- Training Data: 4,700 failures
- Accuracy: 90.1%

### **Model Serving API:**

POST /api/v1/models/predict

POST /api/v1/models/batch-predict

GET /api/v1/models/{model\_name}/metadata

POST /api/v1/models/explain (SHAP values)

### **2.1.4 Master Agent Orchestration Service**

**Purpose:** Central coordinator for all AI agents

## Responsibilities:

- Monitor all data streams and agent status
- Route tasks to appropriate worker agents
- Enforce security policies via UEBA
- Handle escalations and edge cases
- Maintain vehicle state machines
- Generate executive dashboards

## Technical Specifications:

- **Language:** Python (FastAPI with asyncio)
- **Message Queue:** RabbitMQ with dead-letter queues
- **State Storage:** Redis (vehicle state cache)
- **Orchestration Pattern:** Saga pattern for distributed transactions
- **Concurrency:** AsyncIO with worker pool (50-100 concurrent tasks)

## State Machine Per Vehicle:

States:

- HEALTHY (no issues detected)
- MONITORING (anomaly detected, confidence <70%)
- PREDICTION\_CONFIRMED (failure predicted, confidence ≥85%)
- CUSTOMER\_CONTACTED (engagement initiated)
- APPOINTMENT\_SCHEDULED
- SERVICE\_IN\_PROGRESS
- SERVICE\_COMPLETED
- FEEDBACK\_COLLECTED

Transitions:

HEALTHY → MONITORING (anomaly detected)  
MONITORING → PREDICTION\_CONFIRMED (diagnosis confirms issue)  
MONITORING → HEALTHY (false alarm, confidence drops)  
PREDICTION\_CONFIRMED → CUSTOMER\_CONTACTED (engagement agent called)  
CUSTOMER\_CONTACTED → APPOINTMENT\_SCHEDULED (customer confirmed)  
APPOINTMENT\_SCHEDULED → SERVICE\_IN\_PROGRESS (check-in at service center)  
SERVICE\_IN\_PROGRESS → SERVICE\_COMPLETED (work order closed)  
SERVICE\_COMPLETED → FEEDBACK\_COLLECTED (survey submitted)  
FEEDBACK\_COLLECTED → HEALTHY (cycle reset)

## Task Routing Logic:

```
def route_task(prediction_result):  
    if prediction_result['confidence'] >= 0.85:
```



```

# High confidence: Proceed with customer engagement
publish_to_queue('customer_engagement', {
    'vehicle_id': prediction_result['vehicle_id'],
    'component': prediction_result['component'],
    'days_to_failure': prediction_result['days_to_failure'],
    'priority': classify_priority(prediction_result),
    'customer_profile': get_customer_profile(vehicle_id)
})
elif prediction_result['confidence'] >= 0.70:
    # Medium confidence: Enhanced monitoring
    schedule_enhanced_monitoring(prediction_result['vehicle_id'])
else:
    # Low confidence: Log for pattern analysis
    log_low_confidence_prediction(prediction_result)

```

## UEBA Integration:

- Every agent action passes through UEBA check before execution
- Anomaly scores >75 block action and alert security team
- Anomaly scores >90 disable agent immediately

## 2.1.5 Worker Agent Services

### A. Data Analysis Agent

**Purpose:** Continuous monitoring and early warning detection

#### Responsibilities:

- Subscribe to Kafka telemetry stream
- Compute features in real-time
- Detect statistical anomalies
- Publish alerts to diagnosis queue

#### Implementation Details:

- **Deployment:** Kubernetes StatefulSet (for state persistence)
- **Scaling:** Horizontal (partition Kafka topics by VIN hash)
- **Anomaly Detection Methods:**
  - Statistical: Z-score, IQR, Grubbs' test
  - ML-based: Isolation Forest, One-Class SVM
  - Time-series: ARIMA residuals, Prophet anomaly detection

### B. Diagnosis Agent

**Purpose:** Run ML models and rule-based diagnosis

**Responsibilities:**

- Receive anomaly alerts
- Execute ML model inference
- Apply rule-based refinements
- Generate explainable diagnoses
- Classify urgency (P1/P2/P3)

**Rule Engine Examples:**

# Rule 1: Battery voltage override

```
if battery_voltage_24h_mean < 12.0 and cranking_voltage < 9.5:  
    escalate_risk_level('battery', from_level='MEDIUM', to_level='CRITICAL')  
    days_to_failure = 2 # Override model prediction
```

# Rule 2: Brake pad minimum threshold

```
if brake_pad_thickness < 2.5: # Legal minimum in India  
    escalate_risk_level('brake_pads', to_level='CRITICAL')  
    days_to_failure = 0
```

# Rule 3: Model-year specific adjustments

```
if vehicle_model == 'ModelX' and model_year == 2020:  
    # Known issue with transmission in 2020 batch  
    if component == 'transmission':  
        adjust_confidence(multiplier=1.15)
```

**Priority Classification:**

P1 (Critical - 24-48h):

- Safety-related failures (brakes, steering)
- Imminent breakdown (battery dead, transmission failure)
- Vehicle undrivable

P2 (High - 5-10 days):

- Drivability impact (engine misfires, suspension)
- Preventable breakdown (alternator, fuel pump)
- Expensive cascading failures

P3 (Moderate - 10-30 days):

- Preventive maintenance (filters, fluids)
- Comfort features (AC, power windows)
- Schedule at customer convenience

## C. Customer Engagement Agent

**Purpose:** Proactive voice-first customer communication

### Responsibilities:

- Initiate phone calls via Twilio
- Conduct natural conversations using GPT-4
- Handle objections and answer questions
- Provide appointment options
- Fall back to SMS/app if call fails

### Technical Specifications:

- **Primary Channel:** Voice (73% conversion rate)
- **Secondary Channel:** App push notification
- **Tertiary Channel:** SMS with booking link
- **Language Support:** English, Hindi, Regional languages
- **Call Duration Target:** <3 minutes average
- **Retry Logic:** 3 attempts over 48 hours

### Conversation Architecture:

User Input (Speech)

- Whisper ASR (Speech-to-Text)
- Intent Classification (Fine-tuned BERT)
- Context Management (Conversation state)
- GPT-4 Response Generation (with custom system prompt)
- Google Cloud TTS (Text-to-Speech)
- User Output (Speech)

### System Prompt for GPT-4:

You are Maya, a friendly vehicle maintenance assistant for AutoCare. Your goal is to explain predicted vehicle issues in simple terms and persuade customers to book preventive maintenance appointments.

Guidelines:

- Be warm, empathetic, and patient
- Use simple language, avoid technical jargon
- Build trust by referencing their vehicle history
- Handle objections with understanding
- Never be pushy; respect customer autonomy

- If unsure, escalate to human advisor

Customer Context:

- Name: {customer\_name}
- Vehicle: {vehicle\_model} {vehicle\_year}
- Last service: {last\_service\_date}
- Current issue: {predicted\_failure}

Start the conversation with a friendly greeting and permission check.

**Objection Handling Database:** Store common objections with pre-tested response templates:

```
{  
  "objection": "I haven't noticed any problems",  
  "response_template": "That's actually quite common! {component} wear happens gradually, so  
changes aren't noticeable day-to-day. However, our sensors detected {specific_metric} which  
indicates {explanation}. Think of it like a dental checkup - we catch issues before they become  
painful and expensive.",  
  "effectiveness_score": 0.87  
}
```

## D. Scheduling Agent

**Purpose:** Intelligent appointment optimization

**Responsibilities:**

- Query service center availability
- Check parts inventory
- Match technician skills to job requirements
- Optimize for customer location and preferences
- Reserve parts upon confirmation
- Send calendar invites and reminders

**Optimization Algorithm:**

Multi-objective optimization considering:

1. Customer travel distance (minimize)
2. Wait time to appointment (minimize)
3. Parts availability (maximize certainty)
4. Service center load balancing
5. Technician skill match

### Scoring Function:

```
def score_appointment_option(center, slot, customer, job):
    score = (
        weight_distance * distance_score(customer.location, center.location) +
        weight_availability * availability_score(slot, job.urgency) +
        weight_parts * parts_availability_score(center, job.parts_needed) +
        weight_skill * technician_match_score(center, job.complexity) +
        weight_balance * load_balance_score(center.current_utilization)
    )
    return score

weights = {
    'distance': 0.30,
    'availability': 0.25,
    'parts': 0.25,
    'skill': 0.15,
    'balance': 0.05
}
```

### Parts Reservation System:

- Soft hold (15 minutes) while awaiting customer confirmation
- Hard reservation upon booking
- Auto-release if appointment cancelled
- Substitute part lookup if primary part unavailable

### E. Feedback Agent

**Purpose:** Post-service engagement and data loop closure

#### Responsibilities:

- Detect service completion
- Send multi-channel feedback requests
- Collect NPS and detailed feedback
- Perform sentiment analysis
- Update vehicle maintenance history
- Trigger quality insights analysis

#### Feedback Collection Flow:

Service Completed

↓ (wait 2 hours)

SMS Survey (5-star rating)

↓ (if no response after 4 hours)  
App Notification (detailed survey)  
↓ (if no response after 24 hours)  
Voice Call (brief survey)  
↓  
Store feedback + Sentiment Analysis  
↓  
Update vehicle history + Recalibrate predictions

### **Sentiment Analysis:**

- Use fine-tuned BERT for automotive service sentiment
- Classify: Positive / Neutral / Negative
- Extract themes: wait\_time, cost, staff\_behavior, quality, communication
- Flag critical issues (rating  $\leq 2$ ) for immediate escalation

### **F. Quality Insights Module**

**Purpose:** Manufacturing feedback loop and RCA/CAPA generation

#### **Responsibilities:**

- Aggregate field failure patterns
- Correlate with manufacturing data
- Perform statistical root cause analysis
- Generate automated CAPA reports
- Distribute reports to engineering/quality teams
- Track CAPA effectiveness

#### **Failure Pattern Detection:**

```
def detect_failure_patterns(time_window_days=30):
    failures = query_service_records(
        service_type='component_failure',
        date_range=last_n_days(time_window_days)
    )

    # Group by component, model, year
    grouped = failures.groupby(['component', 'model_year', 'model'])

    patterns = []
    for (component, year, model), group in grouped:
        failure_rate = len(group) / total_vehicles(year, model)
        baseline_rate = get_historical_failure_rate(component, year, model)
```

```

# Statistical significance test
z_score = (failure_rate - baseline_rate) / std_error(baseline_rate)

if z_score > 3: # 3-sigma = statistically significant
    patterns.append({
        'component': component,
        'model_year': year,
        'model': model,
        'failure_count': len(group),
        'failure_rate': failure_rate,
        'expected_rate': baseline_rate,
        'statistical_significance': z_score,
        'affected_vins': list(group['vin']),
        'urgency': classify_urgency(z_score, len(group))
    })

return sorted(patterns, key=lambda x: x['statistical_significance'], reverse=True)

```

## **Root Cause Analysis Process:**

### **1. Batch Correlation Analysis:**

- Query manufacturing DB for affected VINs
- Check if failures cluster by supplier batch
- Correlation >70% suggests supplier defect

### **2. Time-Based Analysis:**

- Check if failures from narrow production window
- Suggests process change or environmental issue

### **3. Assembly Line Correlation:**

- Check if specific line/shift has higher failures
- Suggests training or equipment issue

### **4. Material Testing:**

- Pull supplier material certificates
- Check spec compliance
- Identify out-of-spec parameters

### **5. Usage Pattern Analysis:**

- Compare driving patterns of failed vs non-failed
- Normal usage + failures = design/manufacturing issue
- Extreme usage + failures = design margin issue

## CAPA Report Generation:

- Automated report using Jinja2 templates
- Include: problem description, RCA, impact assessment, corrective actions, preventive actions
- Distribute via email + Jira ticket creation
- Track implementation status

## 2.2 Database Design

### 2.2.1 TimescaleDB (Time-Series Data)

**Purpose:** Store high-frequency vehicle telemetry

#### Schema:

-- Hypertable for raw telemetry

```
CREATE TABLE telemetry (  
  time TIMESTAMPTZ NOT NULL,  
  vin TEXT NOT NULL,  
  sensor_type TEXT NOT NULL,  
  sensor_value DOUBLE PRECISION,  
  unit TEXT,  
  quality_flag SMALLINT DEFAULT 0, -- 0=good, 1=interpolated, 2=suspect  
  metadata JSONB  
);
```

```
SELECT create_hypertable('telemetry', 'time',  
  chunk_time_interval => INTERVAL '1 day',  
  if_not_exists => TRUE  
);
```

```
CREATE INDEX idx_telemetry_vin_time ON telemetry (vin, time DESC);  
CREATE INDEX idx_telemetry_sensor_type ON telemetry (sensor_type, time DESC);
```

-- Continuous aggregates for performance

```
CREATE MATERIALIZED VIEW telemetry_1min  
WITH (timescaledb.continuous) AS  
SELECT  
  time_bucket('1 minute', time) AS bucket,  
  vin,  
  sensor_type,  
  AVG(sensor_value) AS avg_value,  
  MAX(sensor_value) AS max_value,  
  MIN(sensor_value) AS min_value,
```



```

        STDDEV(sensor_value) AS std_value,
        COUNT(*) AS sample_count
FROM telemetry
GROUP BY bucket, vin, sensor_type;

CREATE MATERIALIZED VIEW telemetry_1hour
WITH (timescaledb.continuous) AS
SELECT
    time_bucket('1 hour', time) AS bucket,
    vin,
    sensor_type,
    AVG(sensor_value) AS avg_value,
    MAX(sensor_value) AS max_value,
    MIN(sensor_value) AS min_value,
    STDDEV(sensor_value) AS std_value
FROM telemetry
GROUP BY bucket, vin, sensor_type;

-- Retention policy
SELECT add_retention_policy('telemetry', INTERVAL '90 days');
```

### **Engineered Features Table:**

```

CREATE TABLE vehicle_features (
    time TIMESTAMPTZ NOT NULL,
    vin TEXT NOT NULL,

    -- Brake features
    brake_pad_thickness_mm DOUBLE PRECISION,
    brake_wear_rate_mm_per_1000km DOUBLE PRECISION,
    hard_brake_count_24h INTEGER,
    brake_pressure_max_1h DOUBLE PRECISION,

    -- Battery features
    battery_voltage_mean_24h DOUBLE PRECISION,
    cranking_voltage DOUBLE PRECISION,
    charge_acceptance_rate DOUBLE PRECISION,
    deep_discharge_events_7d INTEGER,

    -- Engine features
    coolant_temp_max_24h DOUBLE PRECISION,
    oil_pressure_min_7d DOUBLE PRECISION,
    rpm_variance_1h DOUBLE PRECISION,
```

```

-- Transmission features
gear_shift_smoothness_score DOUBLE PRECISION,
transmission_temp_mean_24h DOUBLE PRECISION,

-- Driving pattern features
urban_driving_pct DOUBLE PRECISION,
avg_speed_24h DOUBLE PRECISION,
aggressive_driving_score DOUBLE PRECISION,

-- Context features
ambient_temp_avg DOUBLE PRECISION,
load_estimate DOUBLE PRECISION,

PRIMARY KEY (time, vin)
);

SELECT create_hypertable('vehicle_features', 'time');
CREATE INDEX idx_features_vin_time ON vehicle_features (vin, time DESC);

```

## 2.2.2 PostgreSQL (Relational Data)

### Users Table:

```

CREATE TABLE users (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  email TEXT UNIQUE NOT NULL,
  phone TEXT UNIQUE NOT NULL,
  name TEXT NOT NULL,
  language_preference TEXT DEFAULT 'en', -- en, hi, ta, te, mr
  timezone TEXT DEFAULT 'Asia/Kolkata',

  -- Communication preferences
  preferred_channel TEXT DEFAULT 'voice', -- voice, sms, app
  consent_voice_call BOOLEAN DEFAULT TRUE,
  consent_sms BOOLEAN DEFAULT TRUE,
  consent_push BOOLEAN DEFAULT TRUE,
  do_not_disturb_start TIME,
  do_not_disturb_end TIME,

  -- Customer tier
  loyalty_tier TEXT DEFAULT 'regular', -- regular, silver, gold, platinum

  -- Location
  address_line1 TEXT,

```

```

address_line2 TEXT,
city TEXT,
state TEXT,
postal_code TEXT,
country TEXT DEFAULT 'IN',
gps_location POINT, -- PostGIS

created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW(),
last_login_at TIMESTAMPTZ,

CONSTRAINT valid_channel CHECK (preferred_channel IN ('voice', 'sms', 'app', 'email'))
);

CREATE INDEX idx_users_phone ON users(phone);
CREATE INDEX idx_users_gps ON users USING GIST(gps_location);

```

### **Vehicles Table:**

```

CREATE TABLE vehicles (
  vin TEXT PRIMARY KEY,
  user_id UUID REFERENCES users(id) ON DELETE SET NULL,

  -- Vehicle details
  make TEXT NOT NULL,
  model TEXT NOT NULL,
  variant TEXT,
  model_year INTEGER NOT NULL,
  manufacturing_date DATE,
  production_batch TEXT,
  production_plant TEXT,
  assembly_line TEXT,

  -- Purchase info
  purchase_date DATE,
  purchase_mileage INTEGER DEFAULT 0,

  -- Current status
  current_mileage INTEGER,
  last_telemetry_at TIMESTAMPTZ,
  health_status TEXT DEFAULT 'HEALTHY', -- HEALTHY, MONITORING, AT_RISK,
  CRITICAL

  -- Registration

```

```

registration_number TEXT UNIQUE,
registration_state TEXT,
insurance_expiry DATE,

created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW(),

CONSTRAINT valid_health_status CHECK (health_status IN ('HEALTHY', 'MONITORING',
'AT_RISK', 'CRITICAL'))
);

CREATE INDEX idx_vehicles_user ON vehicles(user_id);
CREATE INDEX idx_vehicles_model ON vehicles(make, model, model_year);
CREATE INDEX idx_vehicles_batch ON vehicles(production_batch);
CREATE INDEX idx_vehicles_health ON vehicles(health_status);

```

### **Predictions Table:**

```

CREATE TABLE predictions (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  vin TEXT REFERENCES vehicles(vin) ON DELETE CASCADE,
  prediction_time TIMESTAMPTZ NOT NULL DEFAULT NOW(),

  -- Prediction details
  component TEXT NOT NULL, -- brake_pads, battery, engine, transmission, etc.
  failure_type TEXT,
  risk_level TEXT NOT NULL, -- NORMAL, LOW, MEDIUM, HIGH, CRITICAL
  confidence DOUBLE PRECISION NOT NULL, -- 0.0 to 1.0

  -- Timeline
  days_to_failure INTEGER,
  failure_date_estimate DATE,
  confidence_interval_lower INTEGER,
  confidence_interval_upper INTEGER,

  -- Priority
  priority TEXT, -- P1, P2, P3
  safety_critical BOOLEAN DEFAULT FALSE,
  drivability_impact TEXT, -- none, minor, moderate, severe

  -- Recommendation
  recommended_action TEXT,
  estimated_service_duration_minutes INTEGER,
  estimated_cost_inr INTEGER,

```

```

-- Model metadata
model_name TEXT NOT NULL,
model_version TEXT NOT NULL,

-- Explainability
shap_values JSONB, -- Top contributing features
contributing_factors JSONB, -- Human-readable explanations

-- Status
status TEXT DEFAULT 'ACTIVE', -- ACTIVE, ACKNOWLEDGED, SERVICED,
FALSE_POSITIVE
acknowledged_at TIMESTAMPTZ,
serviced_at TIMESTAMPTZ,

created_at TIMESTAMPTZ DEFAULT NOW(),

CONSTRAINT valid_risk CHECK (risk_level IN ('NORMAL', 'LOW', 'MEDIUM', 'HIGH',
'CRITICAL')),
CONSTRAINT valid_priority CHECK (priority IN ('P1', 'P2', 'P3'))
);

CREATE INDEX idx_predictions_vin_time ON predictions(vin, prediction_time DESC);
CREATE INDEX idx_predictions_status ON predictions(status) WHERE status = 'ACTIVE';
CREATE INDEX idx_predictions_component ON predictions(component);

```

### **Service Centers Table:**

```

CREATE TABLE service_centers (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  name TEXT NOT NULL,

  -- Location
  address_line1 TEXT NOT NULL,
  address_line2 TEXT,
  city TEXT NOT NULL,
  state TEXT NOT NULL,
  postal_code TEXT NOT NULL,
  gps_location POINT NOT NULL, -- PostGIS

  -- Contact
  phone TEXT NOT NULL,
  email TEXT,
  manager_name TEXT,

```

```

-- Capacity
total_bays INTEGER NOT NULL,
operating_hours_start TIME DEFAULT '09:00',
operating_hours_end TIME DEFAULT '18:00',
weekend_operating BOOLEAN DEFAULT TRUE,

-- Services offered
services_offered TEXT[], -- {brake_service, battery_replacement, engine_repair, ...}

-- Performance metrics
average_rating DOUBLE PRECISION DEFAULT 0,
total_reviews INTEGER DEFAULT 0,
current_utilization_pct DOUBLE PRECISION DEFAULT 0,

is_active BOOLEAN DEFAULT TRUE,
created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_centers_gps ON service_centers USING GIST(gps_location);
CREATE INDEX idx_centers_city ON service_centers(city, state);

```

### **Parts Inventory Table:**

```

CREATE TABLE parts_inventory (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  service_center_id UUID REFERENCES service_centers(id) ON DELETE CASCADE,

  -- Part details
  part_number TEXT NOT NULL,
  part_name TEXT NOT NULL,
  part_category TEXT, -- brake, battery, engine, transmission, electrical
  compatible_models TEXT[], -- {ModelA, ModelB, ...}

  -- Inventory
  quantity_available INTEGER NOT NULL DEFAULT 0,
  quantity_reserved INTEGER NOT NULL DEFAULT 0,
  minimum_stock_level INTEGER DEFAULT 5,

  -- Supplier
  supplier_name TEXT,
  supplier_batch TEXT,
  unit_cost_inr DOUBLE PRECISION,

```

```

-- Reorder
reorder_point INTEGER,
reorder_quantity INTEGER,
lead_time_days INTEGER,

last_restocked_at TIMESTAMPTZ,
updated_at TIMESTAMPTZ DEFAULT NOW(),

UNIQUE(service_center_id, part_number)
);

CREATE INDEX idx_inventory_center_part ON parts_inventory(service_center_id,
part_number);
CREATE INDEX idx_inventory_low_stock ON parts_inventory(service_center_id)
WHERE quantity_available <= minimum_stock_level;

```

### **Appointments Table:**

```

CREATE TABLE appointments (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  vin TEXT REFERENCES vehicles(vin) ON DELETE CASCADE,
  user_id UUID REFERENCES users(id) ON DELETE CASCADE,
  service_center_id UUID REFERENCES service_centers(id),
  prediction_id UUID REFERENCES predictions(id),

  -- Scheduling
  appointment_time TIMESTAMPTZ NOT NULL,
  estimated_duration_minutes INTEGER NOT NULL,
  appointment_end_time TIMESTAMPTZ GENERATED ALWAYS AS
    (appointment_time + (estimated_duration_minutes || ' minutes')::INTERVAL) STORED,

  -- Service details
  service_type TEXT NOT NULL,
  components_to_service TEXT[],
  parts_required JSONB, -- [{part_number, quantity, reserved}]
  assigned_technician_id UUID,
  estimated_cost_inr INTEGER,

  -- Status tracking
  status TEXT DEFAULT 'SCHEDULED', -- SCHEDULED, CONFIRMED, IN_PROGRESS,
  COMPLETED, CANCELLED, NO_SHOW
  booking_source TEXT, -- voice, app, sms, web
  booking_agent_id TEXT, -- Which agent booked this

```

```

-- Communication tracking
confirmation_sent_at TIMESTAMPTZ,
reminder_sent_at TIMESTAMPTZ,
customer_confirmed_at TIMESTAMPTZ,

-- Completion
checked_in_at TIMESTAMPTZ,
service_started_at TIMESTAMPTZ,
service_completed_at TIMESTAMPTZ,
actual_cost_inr INTEGER,

-- Cancellation
cancelled_at TIMESTAMPTZ,
cancellation_reason TEXT,

created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW(),

CONSTRAINT valid_status CHECK (status IN ('SCHEDULED', 'CONFIRMED',
'IN_PROGRESS', 'COMPLETED', 'CANCELLED', 'NO_SHOW'))
);

CREATE INDEX idx_appointments_vin ON appointments(vin);
CREATE INDEX idx_appointments_user ON appointments(user_id);
CREATE INDEX idx_appointments_center_time ON appointments(service_center_id,
appointment_time);
CREATE INDEX idx_appointments_status ON appointments(status) WHERE status IN
('SCHEDULED', 'CONFIRMED');

```

### **Maintenance History Table:**

```

CREATE TABLE maintenance_history (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  vin TEXT REFERENCES vehicles(vin) ON DELETE CASCADE,
  appointment_id UUID REFERENCES appointments(id),

  -- Service details
  service_date DATE NOT NULL,
  service_type TEXT NOT NULL,
  mileage_at_service INTEGER,

  -- Work performed
  components_serviced TEXT[],

```



```

parts_replaced JSONB, -- [{part_number, part_name, quantity, cost}]
labor_hours DOUBLE PRECISION,

-- Cost breakdown
parts_cost_inr INTEGER,
labor_cost_inr INTEGER,
total_cost_inr INTEGER,

-- Technician
technician_id UUID,
technician_name TEXT,
technician_notes TEXT,

-- Quality
service_center_id UUID REFERENCES service_centers(id),
warranty_applicable BOOLEAN DEFAULT TRUE,
warranty_months INTEGER DEFAULT 6,

-- RCA tags (for quality analysis)
failure_mode TEXT,
root_cause_category TEXT,
capa_id TEXT, -- Link to CAPA report if generated

created_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_maintenance_vin_date ON maintenance_history(vin, service_date DESC);
CREATE INDEX idx_maintenance_components ON maintenance_history USING
GIN(components_serviced);
CREATE INDEX idx_maintenance_capa ON maintenance_history(capa_id) WHERE capa_id IS
NOT NULL;

```

### **Feedback Table:**

```

CREATE TABLE feedback (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),
  appointment_id UUID REFERENCES appointments(id) ON DELETE CASCADE,
  user_id UUID REFERENCES users(id),
  vin TEXT REFERENCES vehicles(vin),

  -- Ratings (1-5 scale)
  overall_rating SMALLINT CHECK (overall_rating BETWEEN 1 AND 5),
  service_quality_rating SMALLINT CHECK (service_quality_rating BETWEEN 1 AND 5),
  staff_behavior_rating SMALLINT CHECK (staff_behavior_rating BETWEEN 1 AND 5),

```

```

wait_time_rating SMALLINT CHECK (wait_time_rating BETWEEN 1 AND 5),
facility_cleanliness_rating SMALLINT CHECK (facility_cleanliness_rating BETWEEN 1 AND
5),
value_for_money_rating SMALLINT CHECK (value_for_money_rating BETWEEN 1 AND 5),

-- NPS
nps_score SMALLINT CHECK (nps_score BETWEEN 0 AND 10),
nps_category TEXT GENERATED ALWAYS AS (
    CASE
        WHEN nps_score >= 9 THEN 'PROMOTER'
        WHEN nps_score >= 7 THEN 'PASSIVE'
        ELSE 'DETRACTOR'
    END
) STORED,

-- Detailed feedback
comments TEXT,
positive_aspects TEXT[],
improvement_areas TEXT[],

-- Sentiment analysis
sentiment TEXT, -- POSITIVE, NEUTRAL, NEGATIVE
sentiment_score DOUBLE PRECISION, -- -1.0 to 1.0
themes_extracted TEXT[], -- [wait_time, cost, staff_behavior, ...]

-- Collection metadata
collection_channel TEXT, -- sms, app, voice
collected_at TIMESTAMPTZ DEFAULT NOW(),

-- Follow-up
requires_follow_up BOOLEAN DEFAULT FALSE,
follow_up_completed BOOLEAN DEFAULT FALSE,
follow_up_notes TEXT,

created_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_feedback_appointment ON feedback(appointment_id);
CREATE INDEX idx_feedback_nps ON feedback(nps_category);
CREATE INDEX idx_feedback_low_rating ON feedback(overall_rating) WHERE overall_rating
<= 3;

```

**Agent Audit Log:**

```

CREATE TABLE agent_audit_log (
  id UUID PRIMARY KEY DEFAULT gen_random_uuid(),

  -- Agent identification
  agent_name TEXT NOT NULL,
  agent_type TEXT NOT NULL, -- master, data_analysis, diagnosis, engagement, scheduling,
  feedback, quality
  agent_instance_id TEXT, -- For distributed deployments

  -- Action details
  action_type TEXT NOT NULL, -- API_CALL, DATABASE_QUERY, MESSAGE_SENT,
  PREDICTION_MADE, etc.
  action_description TEXT,

  -- Context
  vehicle_id TEXT,
  user_id UUID,
  appointment_id UUID,

  -- Request/Response
  request_payload JSONB,
  response_payload JSONB,

  -- Result
  status TEXT NOT NULL, -- SUCCESS, FAILURE, BLOCKED, TIMEOUT
  error_message TEXT,

  -- Performance
  execution_time_ms INTEGER,

  -- Security (UEBA)
  anomaly_score DOUBLE PRECISION,
  anomaly_flags TEXT[], -- [UNUSUAL_TIME, UNAUTHORIZED_TABLE,
  HIGH_FREQUENCY, ...]
  ueba_action_taken TEXT, -- ALLOWED, BLOCKED, THROTTLED, ALERTED

  -- Timestamps
  timestamp TIMESTAMPTZ NOT NULL DEFAULT NOW(),

  -- Retention
  partition_key TEXT GENERATED ALWAYS AS (to_char(timestamp, 'YYYY-MM')) STORED
);

CREATE INDEX idx_audit_agent_time ON agent_audit_log(agent_name, timestamp DESC);

```

```
CREATE INDEX idx_audit_status ON agent_audit_log(status) WHERE status IN ('FAILURE',  
'BLOCKED');  
CREATE INDEX idx_audit_anomaly ON agent_audit_log(anomaly_score DESC) WHERE  
anomaly_score >= 75;
```

-- Partition by month for performance

```
CREATE TABLE agent_audit_log_2025_10 PARTITION OF agent_audit_log  
FOR VALUES IN ('2025-10');
```

### **Manufacturing Data Integration Tables:**

-- Production records (synced from manufacturing ERP)

```
CREATE TABLE production_records (  
  vin TEXT PRIMARY KEY,  
  production_date DATE NOT NULL,  
  production_plant TEXT NOT NULL,  
  assembly_line TEXT NOT NULL,  
  shift TEXT, -- morning, afternoon, night
```

-- Quality control

```
qc_passed BOOLEAN DEFAULT TRUE,  
qc_test_results JSONB,  
qc_inspector_id TEXT,
```

-- Supplier batches (critical for RCA)

```
engine_batch TEXT,  
transmission_batch TEXT,  
brake_system_batch TEXT,  
battery_batch TEXT,  
electrical_system_batch TEXT,
```

-- Environmental conditions during production

```
ambient_temp_celsius DOUBLE PRECISION,  
humidity_pct DOUBLE PRECISION,
```

```
created_at TIMESTAMPTZ DEFAULT NOW(),
```

```
imported_at TIMESTAMPTZ DEFAULT NOW()
```

```
);
```

```
CREATE INDEX idx_production_date ON production_records(production_date);
```

```
CREATE INDEX idx_production_plant_line ON production_records(production_plant,  
assembly_line);
```

```
CREATE INDEX idx_production_brake_batch ON production_records(brake_system_batch);
```

```

-- CAPA records
CREATE TABLE capa_records (
  id TEXT PRIMARY KEY, -- e.g., CAPA-2024-Q4-0047

  -- Problem identification
  component TEXT NOT NULL,
  failure_mode TEXT NOT NULL,
  severity_level TEXT NOT NULL, -- LOW, MEDIUM, HIGH, CRITICAL

  -- Affected vehicles
  affected_vins TEXT[],
  affected_count INTEGER,
  vehicles_at_risk INTEGER,

  -- Root cause
  root_cause_hypothesis TEXT,
  root_cause_confidence TEXT, -- LOW, MEDIUM, HIGH
  supplier_batch TEXT,
  production_window_start DATE,
  production_window_end DATE,

  -- Impact assessment
  customer_impact TEXT,
  financial_impact_inr INTEGER,
  regulatory_impact TEXT,

  -- Actions
  corrective_actions JSONB, -- [{action, responsible, deadline, status}]
  preventive_actions JSONB,

  -- Status tracking
  status TEXT DEFAULT 'OPEN', -- OPEN, IN_PROGRESS, VERIFICATION, CLOSED
  opened_date DATE NOT NULL,
  target_closure_date DATE,
  actual_closure_date DATE,

  -- Approvals
  prepared_by TEXT,
  reviewed_by TEXT,
  approved_by TEXT[],

  -- Effectiveness
  effectiveness_measured BOOLEAN DEFAULT FALSE,
  effectiveness_score DOUBLE PRECISION, -- 0-100

```

```

failure_rate_before DOUBLE PRECISION,
failure_rate_after DOUBLE PRECISION,

created_at TIMESTAMPTZ DEFAULT NOW(),
updated_at TIMESTAMPTZ DEFAULT NOW()
);

CREATE INDEX idx_capa_status ON capa_records(status);
CREATE INDEX idx_capa_severity ON capa_records(severity_level);
CREATE INDEX idx_capa_component ON capa_records(component);

```

## 2.3 API Design

### 2.3.1 Public APIs (External Integrations)

**Base URL:** <https://api.autocare.com/v1>

#### Authentication:

- API Key (for vehicle/telemetry ingestion)
- JWT Bearer Token (for user-facing APIs)
- OAuth 2.0 (for third-party integrations)

#### Rate Limiting:

- Telemetry ingestion: 10,000 requests/minute per API key
- User APIs: 100 requests/minute per user
- Admin APIs: 1,000 requests/minute per account

#### Telemetry Ingestion:

```

POST /telemetry/ingest
Content-Type: application/json
X-API-Key: {api_key}

```

#### Request Body:

```

{
  "vin": "1HGBH41JXMN109186",
  "timestamp": "2025-10-13T10:30:00Z",
  "sensors": [
    {"type": "speed", "value": 65.5, "unit": "km/h"},
    {"type": "rpm", "value": 2800, "unit": "rpm"},
    {"type": "battery_voltage", "value": 12.6, "unit": "V"},
    {"type": "brake_pad_thickness", "value": 4.2, "unit": "mm"},
  ]
}

```

```
{
  "type": "coolant_temp", "value": 92, "unit": "celsius"},
  {"type": "oil_pressure", "value": 35, "unit": "psi"}
],
"gps": {
  "latitude": 12.9716,
  "longitude": 77.5946,
  "altitude": 920
},
"odometer": 18420,
"fuel_level": 45.5,
"dtc_codes": []
}
```

Response (200 OK):

```
{
  "status": "accepted",
  "message_id": "msg_abc123",
  "processed_at": "2025-10-13T10:30:01Z"
}
```

### Batch Telemetry Ingestion:

POST /telemetry/batch  
Content-Type: application/json  
X-API-Key: {api\_key}

Request Body:

```
{
  "vin": "1HGBH41JXMN109186",
  "readings": [
    {
      "timestamp": "2025-10-13T10:30:00Z",
      "sensors": [...]
    },
    {
      "timestamp": "2025-10-13T10:31:00Z",
      "sensors": [...]
    }
  ]
}
```

### Vehicle Registration:

POST /vehicles

Authorization: Bearer {jwt\_token}

Request Body:

```
{
  "vin": "1HGBH41JXMN109186",
  "make": "Honda",
  "model": "City",
  "variant": "VX",
  "model_year": 2023,
  "purchase_date": "2023-03-15",
  "purchase_mileage": 0,
  "registration_number": "KA01AB1234"
}
```

Response (201 Created):

```
{
  "vin": "1HGBH41JXMN109186",
  "status": "registered",
  "health_monitoring_enabled": true,
  "telemetry_api_key": "veh_key_xyz789"
}
```

### Vehicle Health Status:

GET /vehicles/{vin}/health

Authorization: Bearer {jwt\_token}

Response (200 OK):

```
{
  "vin": "1HGBH41JXMN109186",
  "health_status": "AT_RISK",
  "last_updated": "2025-10-13T10:30:00Z",
  "current_mileage": 18420,
  "predictions": [
    {
      "component": "brake_pads",
      "risk_level": "HIGH",
      "confidence": 0.89,
      "days_to_failure": 12,
      "estimated_cost": 2800,
      "recommended_action": "Schedule brake pad replacement within 7 days"
    }
  ],
}
```



```
"upcoming_appointments": [  
  {  
    "appointment_id": "appt_123",  
    "service_center": "Indiranagar Service Center",  
    "appointment_time": "2025-10-20T10:00:00Z"  
  }  
]  
}
```

### **Appointment Booking:**

POST /appointments

Authorization: Bearer {jwt\_token}

Request Body:

```
{  
  "vin": "1HGBH41JXMN109186",  
  "service_type": "brake_pad_replacement",  
  "preferred_date": "2025-10-20",  
  "preferred_time_slot": "morning", // morning, afternoon, evening  
  "notes": "Customer requested weekend appointment"  
}
```

Response (201 Created):

```
{  
  "appointment_id": "appt_123",  
  
  "appointment_id": "appt_123",  
  "vin": "1HGBH41JXMN109186",  
  "service_center": {  
    "id": "center_001",  
    "name": "Indiranagar Service Center",  
    "address": "123 Main Road, Indiranagar, Bangalore 560038",  
    "phone": "+91-80-12345678",  
    "distance_km": 4.2  
  },  
  "appointment_time": "2025-10-20T10:00:00Z",  
  "estimated_duration_minutes": 90,  
  "estimated_cost": 2800,  
  "parts_reserved": true,  
  "status": "CONFIRMED",  
  "calendar_invite_sent": true  
}
```

## Trip Planning / Pre-Trip Check:

POST /vehicles/{vin}/trip-check  
Authorization: Bearer {jwt\_token}

Request Body:

```
{
  "origin": {
    "latitude": 12.9716,
    "longitude": 77.5946
  },
  "destination": {
    "latitude": 15.2993,
    "longitude": 74.1240
  },
  "departure_date": "2025-10-25",
  "estimated_duration_hours": 8,
  "passenger_count": 4
}
```

Response (200 OK):

```
{
  "trip_safe": false,
  "recommendations": [
    {
      "priority": "HIGH",
      "component": "brake_pads",
      "issue": "Current brake pad thickness is 2.8mm. For a 500km highway trip with 4 passengers, we recommend replacement to ensure safety.",
      "estimated_wear_during_trip": "0.6mm",
      "risk_level": "HIGH",
      "action_required": "Replace brake pads before trip"
    }
  ],
  "suggested_services": [
    {
      "service": "brake_pad_replacement",
      "urgency": "before_trip",
      "estimated_cost": 2800
    },
    {
      "service": "tire_pressure_check",
      "urgency": "recommended",
      "estimated_cost": 0
    }
  ]
}
```

```
}  
],  
"nearby_service_centers": [  
  {  
    "id": "center_001",  
    "name": "Indiranagar Service Center",  
    "distance_km": 4.2,  
    "available_slots": ["2025-10-23T09:00:00Z", "2025-10-23T14:00:00Z"]  
  }  
]  
}
```

### Feedback Submission:

POST /appointments/{appointment\_id}/feedback

Authorization: Bearer {jwt\_token}

Request Body:

```
{  
  "overall_rating": 5,  
  "service_quality_rating": 5,  
  "staff_behavior_rating": 5,  
  "wait_time_rating": 4,  
  "facility_cleanliness_rating": 5,  
  "nps_score": 9,  
  "comments": "Excellent service! The technician explained everything clearly.",  
  "would_recommend": true  
}
```

Response (201 Created):

```
{  
  "feedback_id": "fb_456",  
  "status": "received",  
  "thank_you_points_awarded": 100  
}
```

## 2.3.2 Internal APIs (Agent Communication)

### Master Agent Orchestration:

POST /internal/agents/task

X-Agent-Auth: {internal\_token}

Request Body:

```
{
  "task_id": "task_abc123",
  "task_type": "CUSTOMER_ENGAGEMENT",
  "priority": "HIGH",
  "payload": {
    "vehicle_id": "1HGBH41JXMN109186",
    "prediction": {
      "component": "brake_pads",
      "risk_level": "HIGH",
      "confidence": 0.89,
      "days_to_failure": 12
    },
    "customer_profile": {
      "name": "Rajesh Sharma",
      "phone": "+91-9876543210",
      "language": "en",
      "preferred_channel": "voice"
    }
  },
  "timeout_seconds": 300
}
```

Response (202 Accepted):

```
{
  "task_id": "task_abc123",
  "status": "QUEUED",
  "estimated_start_time": "2025-10-13T10:35:00Z"
}
```

### **Model Inference:**

POST /internal/models/predict  
X-Agent-Auth: {internal\_token}

Request Body:

```
{
  "vin": "1HGBH41JXMN109186",
  "features": {
    "brake_pad_thickness_mm": 3.2,
    "brake_wear_rate_mm_per_1000km": 0.8,
    "hard_brake_count_24h": 15,
    "urban_driving_pct": 75,
    "vehicle_age_months": 18,

```

```
    "avg_speed_24h": 32.5
  },
  "model_name": "brake_failure_predictor",
  "explain": true
}
```

Response (200 OK):

```
{
  "prediction": {
    "failure_probability": 0.89,
    "risk_class": "HIGH",
    "days_to_failure": 12,
    "confidence_interval": [9, 15]
  },
  "explanation": {
    "shap_values": [
      {"feature": "brake_wear_rate_mm_per_1000km", "contribution": 0.34},
      {"feature": "brake_pad_thickness_mm", "contribution": 0.28},
      {"feature": "hard_brake_count_24h", "contribution": 0.19}
    ],
    "human_readable": "High wear rate (0.8mm/1000km) and frequent hard braking are the primary risk factors."
  },
  "model_version": "v2.3.1",
  "inference_time_ms": 42
}
```

### UEBA Check:

POST /internal/ueba/validate  
X-Agent-Auth: {internal\_token}

Request Body:

```
{
  "agent_name": "scheduling_agent_01",
  "action_type": "DATABASE_QUERY",
  "action_details": {
    "query_type": "SELECT",
    "table": "service_centers",
    "row_count": 3,
    "payload_size_bytes": 1024
  },
  "context": {
    "timestamp": "2025-10-13T10:30:00Z",
```

```
"source_ip": "10.0.1.45"
}
}
```

Response (200 OK):

```
{
  "allowed": true,
  "anomaly_score": 12,
  "risk_level": "LOW",
  "action": "ALLOWED",
  "baseline_deviation": 0.3
}
```

## 2.4 Message Queue Architecture

### Kafka Topics:

topics:

```
# Raw telemetry ingestion
- name: vehicle.telemetry.raw
  partitions: 50
  replication_factor: 3
  retention_ms: 604800000 # 7 days
```

```
# Processed features
- name: vehicle.features.computed
  partitions: 50
  replication_factor: 3
  retention_ms: 2592000000 # 30 days
```

```
# Anomaly alerts
- name: vehicle.anomalies.detected
  partitions: 20
  replication_factor: 3
  retention_ms: 2592000000 # 30 days
```

```
# Predictions
- name: vehicle.predictions.generated
  partitions: 20
  replication_factor: 3
  retention_ms: 7776000000 # 90 days
```

```
# Agent tasks (RabbitMQ handles better)
```

- # - customer\_engagement\_tasks
- # - scheduling\_tasks
- # - feedback\_tasks

## **RabbitMQ Queues:**

queues:

- # Worker agent task queues
  - name: tasks.customer\_engagement
    - durable: true
    - message\_ttl: 3600000 # 1 hour
    - max\_priority: 10
    - dead\_letter\_exchange: dlx.tasks
  - name: tasks.scheduling
    - durable: true
    - message\_ttl: 1800000 # 30 minutes
    - max\_priority: 10
    - dead\_letter\_exchange: dlx.tasks
  - name: tasks.feedback
    - durable: true
    - message\_ttl: 86400000 # 24 hours
    - max\_priority: 5
    - dead\_letter\_exchange: dlx.tasks
  - name: tasks.quality\_insights
    - durable: true
    - message\_ttl: 604800000 # 7 days
    - max\_priority: 5
    - dead\_letter\_exchange: dlx.tasks
- # Dead letter queues
  - name: dlx.tasks.customer\_engagement
    - durable: true

# Priority routing

exchanges:

- name: ex.tasks
  - type: topic
  - durable: true
  - routing\_keys:
    - "task.engagement.#"
    - "task.scheduling.#"

- "task.feedback.##"

## 2.5 Caching Strategy

### Redis Cache Layers:

cache\_layers:

# Hot vehicle state (very frequently accessed)

- name: vehicle\_state

prefix: "veh:{vin}:state"

ttl: 3600 # 1 hour

eviction: LRU

example\_key: "veh:1HGBH41JXMN109186:state"

example\_value:

last\_telemetry: "2025-10-13T10:30:00Z"

health\_status: "AT\_RISK"

active\_predictions: [{component: "brake\_pads", confidence: 0.89}]

scheduled\_appointment: "2025-10-20T10:00:00Z"

# Recent predictions

- name: predictions\_cache

prefix: "pred:{vin}:{component}"

ttl: 7200 # 2 hours

eviction: LRU

# Customer profiles

- name: customer\_profiles

prefix: "user:{user\_id}:profile"

ttl: 1800 # 30 minutes

eviction: LRU

# Service center availability

- name: center\_availability

prefix: "center:{center\_id}:slots:{date}"

ttl: 300 # 5 minutes (frequent updates)

eviction: LRU

# Parts inventory

- name: parts\_inventory

prefix: "inventory:{center\_id}:{part\_number}"

ttl: 600 # 10 minutes

eviction: LRU



```
# Model inference results (short-lived)
- name: inference_cache
  prefix: "inference:{model}:{feature_hash}"
  ttl: 60 # 1 minute
  eviction: LFU
```

```
# Rate limiting
- name: rate_limits
  prefix: "ratelimit:{endpoint}:{api_key}"
  ttl: 60 # 1 minute sliding window
  eviction: None (don't evict)
```

### Cache Invalidation Patterns:

1. **Write-Through:** Updates written to DB first, then cache updated
2. **Cache-Aside:** Read from cache, if miss, read from DB and populate cache
3. **Event-Based:** Kafka events trigger cache invalidation
4. **TTL-Based:** Natural expiration for less critical data

## 2.6 Security Implementation

### 2.6.1 UEBA (User and Entity Behavior Analytics)

#### Baseline Learning Process (30-day training):

For each agent, collect:

- API call patterns (frequency, endpoints, payload sizes)
- Database access patterns (tables, query types, row counts)
- Execution patterns (duration, success rate, error types)
- Communication patterns (message queue activity)
- Temporal patterns (time-of-day, day-of-week)

#### Anomaly Detection Algorithm:

```
def calculate_anomaly_score(agent_id, current_behavior):
    baseline = get_agent_baseline(agent_id)
    score = 0
    violations = []

    # 1. API Call Anomalies (30%)
    for api, calls in current_behavior.api_calls.items():
        if api not in baseline.allowed_apis:
            score += 40
```

```

        violations.append(f"UNAUTHORIZED_API: {api}")
    elif calls.frequency > baseline.mean + (3 * baseline.std):
        score += 15
        violations.append(f"HIGH_FREQUENCY: {api}")

# 2. Data Access Anomalies (30%)
for table in current_behavior.tables_accessed:
    if table in baseline.forbidden_tables:
        score += 50
        violations.append(f"FORBIDDEN_TABLE: {table}")
    if current_behavior.row_count[table] > baseline.typical_rows * 5:
        score += 20
        violations.append(f"EXCESSIVE_DATA_ACCESS: {table}")

# 3. Temporal Anomalies (15%)
if is_unusual_time(current_behavior.timestamp, baseline.activity_hours):
    score += 10
    violations.append("UNUSUAL_TIME")

# 4. Execution Anomalies (15%)
if current_behavior.duration > baseline.mean_duration + (4 * baseline.std_duration):
    score += 10
    violations.append("SLOW_EXECUTION")
if current_behavior.error_rate > baseline.error_rate * 3:
    score += 10
    violations.append("HIGH_ERROR_RATE")

# 5. Communication Anomalies (10%)
if current_behavior.message_count > baseline.mean_messages + (5 *
baseline.std_messages):
    score += 15
    violations.append("EXCESSIVE_MESSAGING")

return {
    'anomaly_score': min(score, 100),
    'violations': violations,
    'severity': classify_severity(score)
}

```

### **Automated Response Actions:**

```

def handle_anomaly(anomaly_result):
    score = anomaly_result['anomaly_score']

```

```

if score >= 90: # CRITICAL
    disable_agent(agent_id)
    revoke_api_credentials(agent_id)
    send_pagerduty_alert()
    preserve_forensic_evidence(agent_id)

elif score >= 75: # HIGH
    block_current_action(agent_id)
    send_security_alert()
    require_manual_review()

elif score >= 50: # MEDIUM
    log_warning()
    increase_monitoring_frequency(from_5min_to_1min)
    send_ops_notification()

else: # LOW
    log_info()

```

## 2.6.2 Data Encryption

### At Rest:

- AES-256 encryption for all databases
- Encrypted EBS volumes for Kubernetes persistent volumes
- S3 server-side encryption (SSE-S3 or SSE-KMS)
- Encrypted Redis with AUTH

### In Transit:

- TLS 1.3 for all API communications
- mTLS for inter-service communication
- WSS (WebSocket Secure) for real-time updates
- MQTT over TLS for vehicle telemetry

### PII Protection:

- Tokenization for phone numbers and emails
- Hash VINs in logs
- Field-level encryption for sensitive customer data
- Data masking in non-production environments

## 2.6.3 Authentication & Authorization

### User Authentication:

- JWT tokens with 1-hour expiration
- Refresh tokens with 30-day expiration
- OAuth 2.0 for third-party integrations
- MFA for admin users

#### **Service-to-Service:**

- mTLS certificates
- Service accounts with scoped permissions
- API keys with IP whitelisting
- Internal token validation via shared secret

#### **RBAC Roles:**

roles:

- name: customer

permissions:

- read:own\_vehicles
- write:own\_appointments
- write:own\_feedback

- name: service\_advisor

permissions:

- read:assigned\_appointments
- write:service\_records
- read:vehicle\_diagnostics

- name: service\_center\_manager

permissions:

- read:center\_appointments
- write:center\_inventory
- read:center\_analytics

- name: engineer

permissions:

- read:all\_predictions
- read:manufacturing\_data
- write:capa\_records

- name: admin

permissions:

- "\*".\*"

---

## 3. FRONTEND ARCHITECTURE

### 3.1 Technology Stack

#### Core Framework:

- Next.js 14 (App Router) with React Server Components
- React 18 with TypeScript
- Server-side rendering for SEO and performance

#### 3D Visualization:

- React Three Fiber (R3F) for 3D rendering
- Three.js r160+ as the underlying engine
- @react-three/drei for helper components
- @react-three/postprocessing for effects (bloom, DOF)

#### UI Components:

- Tailwind CSS for styling
- shadcn/ui for component library
- Framer Motion for animations
- Lucide React for icons

#### Data Visualization:

- ECharts for complex charts
- Recharts for simple charts
- D3.js for custom visualizations

#### State Management:

- Zustand for global state
- React Query (TanStack Query) for server state
- Context API for theme/auth

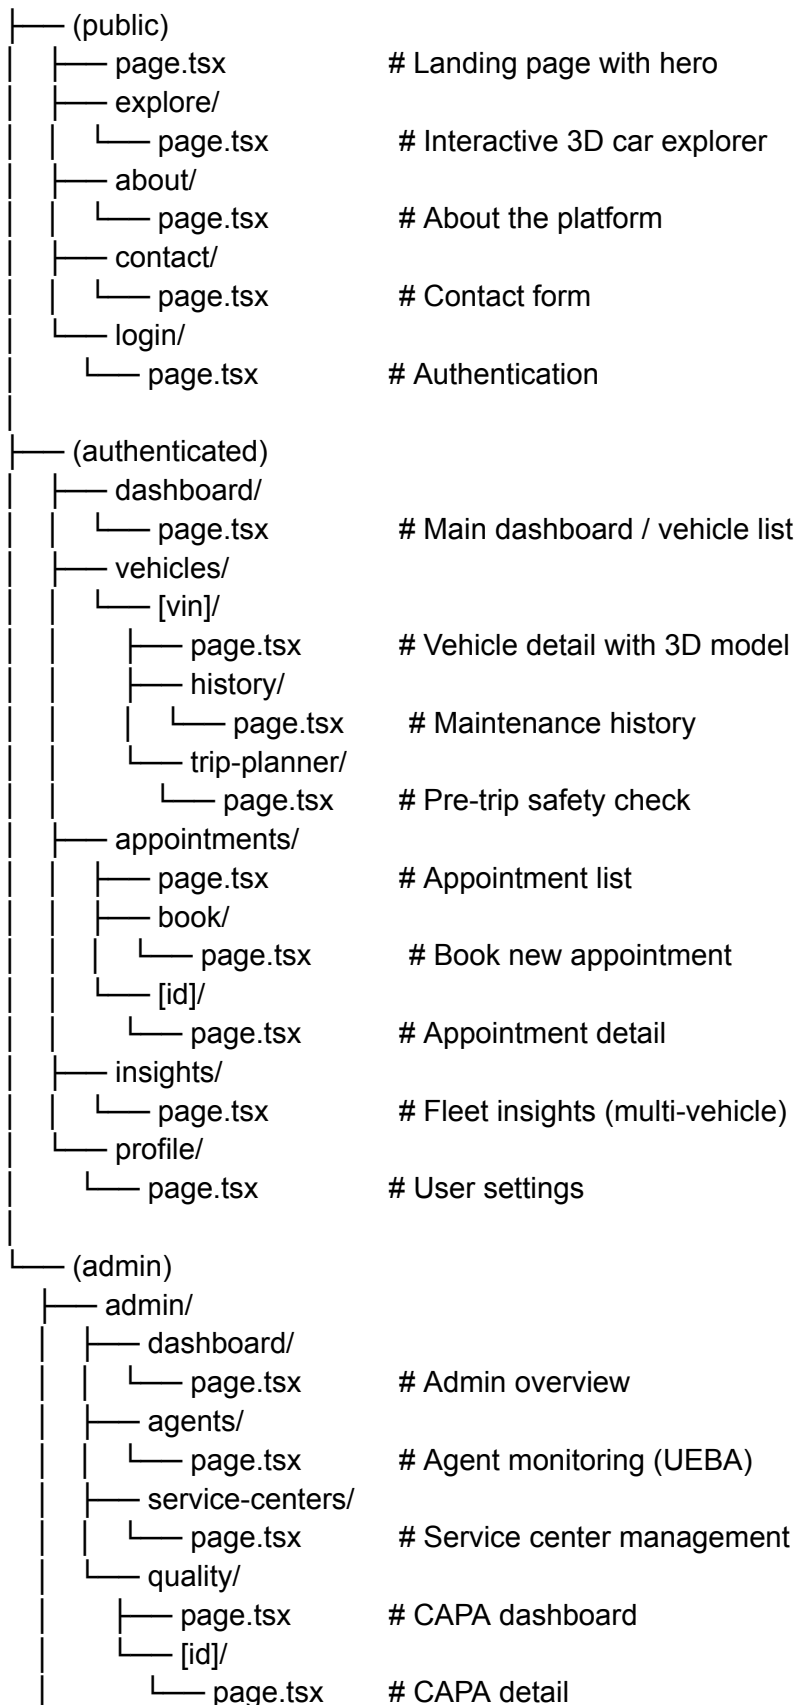
#### Real-Time Communication:

- Socket.io for live updates
- Server-Sent Events for notifications

### 3.2 Site Structure & Pages

#### 3.2.1 Page Architecture

/



## 3.3 Key Components & Features

### 3.3.1 Landing Page

#### Hero Section with 3D Car:

```
// components/Hero3D.tsx
'use client'

import { Canvas } from '@react-three/fiber'
import { OrbitControls, Environment, useGLTF } from '@react-three/drei'
import { Suspense } from 'react'

export function Hero3D() {
  return (
    <div className="h-screen w-full">
      <Canvas camera={{ position: [5, 2, 5], fov: 50 }}>
        <Suspense fallback={<LoadingSpinner />}>
          <Environment preset="sunset" />
          <CarModel />
          <OrbitControls
            enableZoom={false}
            enablePan={false}
            maxPolarAngle={Math.PI / 2}
            minPolarAngle={Math.PI / 3}
          />
        </Suspense>
      </Canvas>

      <div className="absolute inset-0 flex items-center justify-center pointer-events-none">
        <div className="text-center pointer-events-auto">
          <h1 className="text-6xl font-bold mb-4">
            Never Break Down Again
          </h1>
          <p className="text-xl mb-8">
            AI-powered predictive maintenance for your vehicle
          </p>
          <button className="px-8 py-3 bg-blue-600 rounded-lg">
            Get Started
          </button>
        </div>
      </div>
    </div>
  )
}
```

)  
}

### Value Proposition Cards:

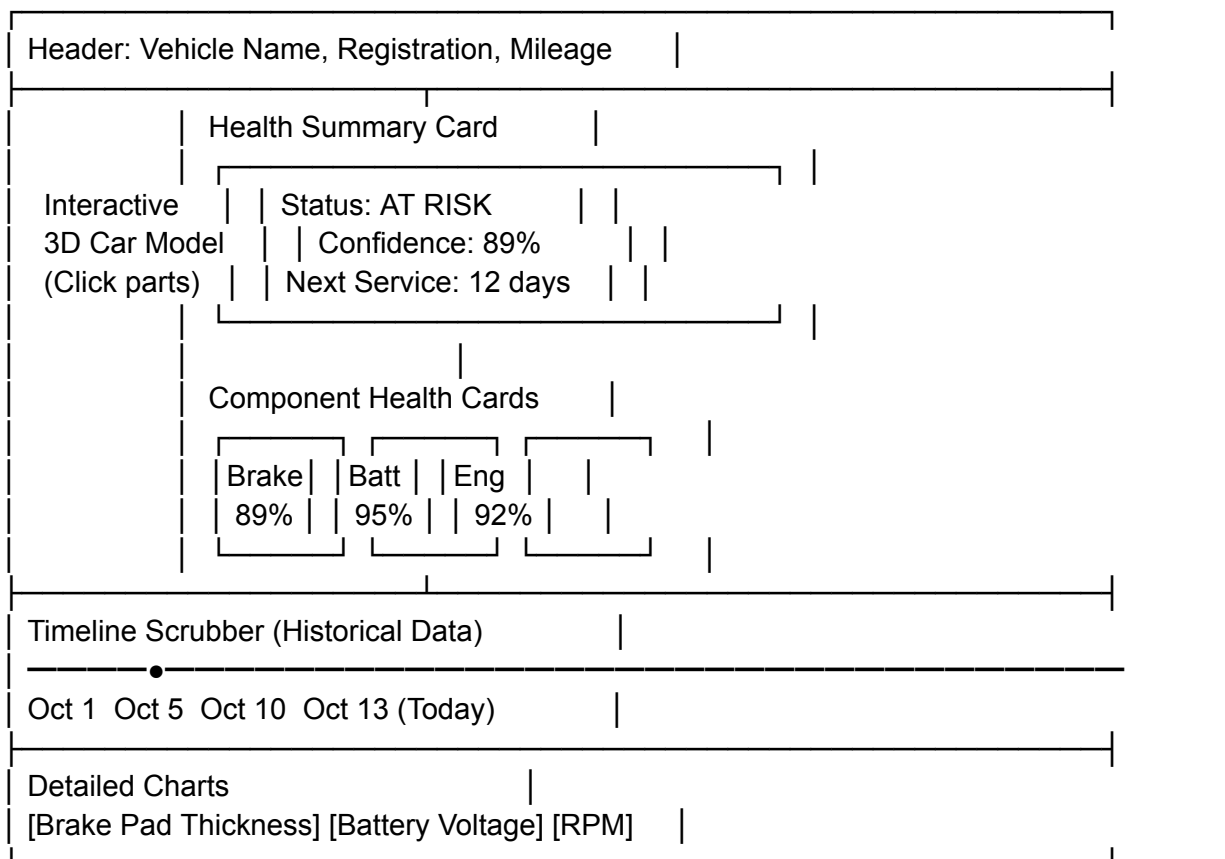
- Real-time monitoring with live sensor visualization
- 93% reduction in breakdowns (animated counter)
- Autonomous scheduling (chatbot demo)
- Manufacturing quality feedback (infographic)

### Social Proof:

- Customer testimonials with photos
- Service center partner logos
- Statistics with animated counters

### 3.3.2 Vehicle Detail Page (Core UX)

#### Layout:





### 3D Car Interaction:

```
// components/InteractiveCar.tsx
'use client'

import { Canvas } from '@react-three/fiber'
import { useGLTF, Html } from '@react-three/drei'
import { useState } from 'react'

function CarModel({ onPartClick, highlightedPart }) {
  const { scene, nodes } = useGLTF('/models/car.glb')

  const parts = {
    brake_front_left: nodes.BrakeFrontLeft,
    brake_front_right: nodes.BrakeFrontRight,
    battery: nodes.Battery,
    engine: nodes.Engine,
    transmission: nodes.Transmission
  }

  return (
    <group>
      {Object.entries(parts).map(([partName, mesh]) => (
        <mesh
          key={partName}
          geometry={mesh.geometry}
          material={mesh.material.clone()}
          onClick={() => onPartClick(partName)}
          onPointerOver={(e) => {
            e.stopPropagation()
            document.body.style.cursor = 'pointer'
            // Glow effect
            e.object.material.emissive.set(0x4444ff)
            e.object.material.emissiveIntensity = 0.3
          }}
          onPointerOut={(e) => {
            document.body.style.cursor = 'default'
            e.object.material.emissive.set(0x000000)
          }}
        >
          {highlightedPart === partName && (
            <Html position={[0, 0.5, 0]}>
              <div className="bg-black/80 text-white px-3 py-2 rounded">
                Click for details
              </div>
            </Html>
          )}
        </mesh>
      ))}
    </group>
  )
}
```

```

        </div>
      </Html>
    )}
  </mesh>
)}
</group>
)
}

```

```

export function InteractiveCar({ vehicleData }) {
  const [selectedPart, setSelectedPart] = useState(null)

  return (
    <div className="relative h-full">
      <Canvas>
        <CarModel
          onPartClick={setSelectedPart}
          highlightedPart={selectedPart}
        />
      </Canvas>

      {selectedPart && (
        <PartDetailPanel
          part={selectedPart}
          data={vehicleData.predictions[selectedPart]}
          onClose={() => setSelectedPart(null)}
        />
      )}
    </div>
  )
}

```

### Component Health Cards:

```

// components/ComponentHealthCard.tsx
interface ComponentHealth {
  component: string
  health_score: number // 0-100
  status: 'healthy' | 'warning' | 'critical'
  last_service: string
  predicted_failure_days: number | null
}

```

```

export function ComponentHealthCard({ component }: { component: ComponentHealth }) {

```

```

const statusColors = {
  healthy: 'bg-green-500',
  warning: 'bg-yellow-500',
  critical: 'bg-red-500'
}

return (
  <div className="border rounded-lg p-4 hover:shadow-lg transition">
    <div className="flex items-center justify-between mb-3">
      <h3 className="font-semibold">{component.component}</h3>
      <span className={`h-3 w-3 rounded-full ${statusColors[component.status]} ` />
    </div>

    <div className="mb-3">
      <div className="flex justify-between text-sm mb-1">
        <span>Health</span>
        <span>{component.health_score}%</span>
      </div>
      <div className="h-2 bg-gray-200 rounded-full overflow-hidden">
        <div
          className={`h-full ${statusColors[component.status]} `}
          style={{ width: `${component.health_score}%` }}
        />
      </div>
    </div>

    {component.predicted_failure_days && (
      <div className="text-sm text-gray-600">
        Predicted service in {component.predicted_failure_days} days
      </div>
    )}

    <div className="text-xs text-gray-500 mt-2">
      Last serviced: {new Date(component.last_service).toLocaleDateString()}
    </div>
  </div>
)
}

```

### Timeline Scrubber:

```

// components/TimelineScrubber.tsx
'use client'

```

```

import { useState, useEffect } from 'react'
import { Line } from 'react-chartjs-2'

export function TimelineScrubber({ vehicleData, onChange }) {
  const [selectedDate, setSelectedDate] = useState(new Date())

  // Historical data from last 90 days
  const historicalData = vehicleData.historical_health_scores

  return (
    <div className="p-4 bg-gray-50 rounded-lg">
      <div className="mb-4">
        <h3 className="text-sm font-semibold mb-2">Health History</h3>
        <Line
          data={{
            labels: historicalData.map(d => d.date),
            datasets: [
              {
                label: 'Overall Health',
                data: historicalData.map(d => d.health_score),
                borderColor: 'rgb(59, 130, 246)',
                tension: 0.4
              }
            ]
          }}
          options={{
            responsive: true,
            interaction: {
              intersect: false,
              mode: 'index'
            },
            plugins: {
              tooltip: {
                callbacks: {
                  afterLabel: (context) => {
                    const date = historicalData[context.dataIndex].date
                    return `Click to view details for ${date}`
                  }
                }
              }
            }
          },
          onClick: (event, elements) => {
            if (elements.length > 0) {
              const index = elements[0].index

```

```

        onChange(historicalData[index].date)
      }
    }
  }}
  />
</div>

<input
  type="range"
  min={0}
  max={historicalData.length - 1}
  value={historicalData.findIndex(d => d.date === selectedDate)}
  onChange={(e) => {
    const date = historicalData[parseInt(e.target.value)].date
    setSelectedDate(date)
    onChange(date)
  }}
  className="w-full"
/>
</div>
)
}

```

### 3.3.3 Trip Planner Page

#### Features:

- Route input (origin → destination)
- Estimated distance and duration
- Pre-trip safety check with component-by-component analysis
- Recommended services before departure
- Nearby service center map

```

// app/vehicles/[vin]/trip-planner/page.tsx
'use client'

```

```

import { useState } from 'react'
import { MapPin, Calendar, Users, AlertTriangle } from 'lucide-react'

```

```

export default function TripPlanner({ params }) {
  const [trip, setTrip] = useState({
    origin: "",
    destination: "",
    date: "",

```

```
    passengers: 1
  })
```

```
const [analysis, setAnalysis] = useState(null)
```

```
const analyzeTripSafety = async () => {
  const response = await fetch(`/api/vehicles/${params.vin}/trip-check`, {
    method: 'POST',
    body: JSON.stringify(trip)
  })
  const data = await response.json()
  setAnalysis(data)
}
```

```
return (
  <div className="max-w-4xl mx-auto p-6">
    <h1 className="text-3xl font-bold mb-6">Plan Your Trip</h1>

    <div className="bg-white rounded-lg shadow p-6 mb-6">
      <div className="grid grid-cols-2 gap-4">
        <div>
          <label className="block text-sm font-medium mb-2">
            <MapPin className="inline w-4 h-4 mr-1" />
            From
          </label>
          <input
            type="text"
            placeholder="Bangalore"
            className="w-full border rounded px-3 py-2"
            value={trip.origin}
            onChange={(e) => setTrip({...trip, origin: e.target.value})}
          />
        </div>

        <div>
          <label className="block text-sm font-medium mb-2">
            <MapPin className="inline w-4 h-4 mr-1" />
            To
          </label>
          <input
            type="text"
            placeholder="Goa"
            className="w-full border rounded px-3 py-2"
            value={trip.destination}
          />
        </div>
      </div>
    </div>
  </div>
)
```

```

        onChange={(e) => setTrip({...trip, destination: e.target.value})}
      />
    </div>

    <div>
      <label className="block text-sm font-medium mb-2">
        <Calendar className="inline w-4 h-4 mr-1" />
        Departure Date
      </label>
      <input
        type="date"
        className="w-full border rounded px-3 py-2"
        value={trip.date}
        onChange={(e) => setTrip({...trip, date: e.target.value})}
      />
    </div>

    <div>
      <label className="block text-sm font-medium mb-2">
        <Users className="inline w-4 h-4 mr-1" />
        Passengers
      </label>
      <input
        type="number"
        min={1}
        max={7}
        className="w-full border rounded px-3 py-2"
        value={trip.passengers}
        onChange={(e) => setTrip({...trip, passengers: parseInt(e.target.value)}}
      />
    </div>
  </div>

  <button
    onClick={analyzeTripSafety}
    className="mt-4 w-full bg-blue-600 text-white py-3 rounded-lg font-semibold
    hover:bg-blue-700"
  >
    Check Trip Safety
  </button>
</div>

{analysis && (
  <div className="space-y-4">

```

```

    {!analysis.trip_safe && (
      <div className="bg-red-50 border-l-4 border-red-500 p-4 rounded">
        <div className="flex items-center mb-2">
          <AlertTriangle className="w-5 h-5 text-red-500 mr-2" />
          <h3 className="font-semibold text-red-800">Action Required Before Trip</h3>
        </div>
        <p className="text-sm text-red-700">
          We've identified issues that should be addressed for a safe journey.
        </p>
      </div>
    )}

    <div className="bg-white rounded-lg shadow p-6">
      <h3 className="font-semibold mb-4">Recommended Services</h3>
      {analysis.recommendations.map((rec, idx) => (
        <div key={idx} className="border-l-4 border-yellow-500 pl-4 mb-4">
          <div className="flex items-center justify-between">
            <div>
              <h4 className="font-medium">{rec.component}</h4>
              <p className="text-sm text-gray-600">{rec.issue}</p>
            </div>
            <span className={`px-3 py-1 rounded text-sm ${
              rec.priority === 'HIGH' ? 'bg-red-100 text-red-800' : 'bg-yellow-100 text-yellow-800'
            }`}>
              {rec.priority}
            </span>
          </div>
        </div>
      ))}
    </div>

    <div className="bg-white rounded-lg shadow p-6">
      <h3 className="font-semibold mb-4">Nearby Service Centers</h3>
      { /* Map component showing service centers */ }
      { /* List of centers with available slots */ }
    </div>
  </div>
)
}

```

### 3.3.4 Dashboard Page



### Fleet Overview (for multi-vehicle users):

```
// app/dashboard/page.tsx
export default async function Dashboard() {
  const vehicles = await getUserVehicles()

  return (
    <div className="p-6">
      <div className="grid grid-cols-1 md:grid-cols-3 gap-6 mb-8">
        <StatCard
          title="Vehicles Monitored"
          value={vehicles.length}
          icon={<Car />}
        />
        <StatCard
          title="Upcoming Services"
          value={vehicles.filter(v => v.next_service_days <= 15).length}
          icon={<Wrench />}
        />
        <StatCard
          title="Health Score"
          value={`$${Math.round(vehicles.reduce((acc, v) => acc + v.health_score, 0) /
vehicles.length)}%`}
          icon={<Heart />}
        />
      </div>

      <div className="grid grid-cols-1 lg:grid-cols-2 gap-6">
        {vehicles.map(vehicle => (
          <VehicleCard key={vehicle.vin} vehicle={vehicle} />
        ))}
      </div>
    </div>
  )
}
```

### 3.3.5 Admin Dashboard

#### Agent Monitoring (UEBA):

```
// app/admin/agents/page.tsx
'use client'

import { useEffect, useState } from 'react'
```

```

import { AlertTriangle, CheckCircle, XCircle } from 'lucide-react'

export default function AgentMonitoring() {
  const [agents, setAgents] = useState([])
  const [alerts, setAlerts] = useState([])

  useEffect(() => {
    // Real-time updates via WebSocket
    const ws = new WebSocket('wss://api.autocare.com/admin/agents/stream')

    ws.onmessage = (event) => {
      const data = JSON.parse(event.data)
      if (data.type === 'AGENT_STATUS') {
        setAgents(prev => updateAgentStatus(prev, data.payload))
      } else if (data.type === 'UEBA_ALERT') {
        setAlerts(prev => [data.payload, ...prev].slice(0, 50))
      }
    }
  })

  return () => ws.close()
}, [])

return (
  <div className="p-6">
    <h1 className="text-3xl font-bold mb-6">Agent Monitoring</h1>

    {/* Live alerts */}
    {alerts.length > 0 && (
      <div className="bg-red-50 border-l-4 border-red-500 p-4 mb-6 rounded">
        <h3 className="font-semibold text-red-800 mb-2">
          <AlertTriangle className="inline w-5 h-5 mr-2" />
          Recent UEBA Alerts
        </h3>
        {alerts.slice(0, 3).map(alert => (
          <div key={alert.id} className="text-sm text-red-700 mb-1">
            {alert.timestamp} - {alert.agent_name}: {alert.violation}
          </div>
        ))}
      </div>
    )}

    {/* Agent status grid */}
    <div className="grid grid-cols-1 md:grid-cols-2 lg:grid-cols-3 gap-6">
      {agents.map(agent => (

```

```

<div key={agent.id} className="bg-white rounded-lg shadow p-6">
  <div className="flex items-center justify-between mb-4">
    <h3 className="font-semibold">{agent.name}</h3>
    {agent.status === 'ACTIVE' ? (
      <CheckCircle className="w-5 h-5 text-green-500" />
    ) : (
      <XCircle className="w-5 h-5 text-red-500" />
    )}
  </div>

  <div className="space-y-2 text-sm">
    <div className="flex justify-between">
      <span className="text-gray-600">Tasks Processed</span>
      <span className="font-medium">{agent.tasks_today}</span>
    </div>
    <div className="flex justify-between">
      <span className="text-gray-600">Success Rate</span>
      <span className="font-medium">{agent.success_rate}%</span>
    </div>
    <div className="flex justify-between">
      <span className="text-gray-600">Anomaly Score</span>
      <span className={`font-medium ${
        agent.anomaly_score > 75 ? 'text-red-600' : 'text-green-600'
      }`}>
        {agent.anomaly_score}/100
      </span>
    </div>
  </div>

  <div className="mt-4 pt-4 border-t">
    <div className="text-xs text-gray-500">
      Last active: {agent.last_active}
    </div>
  </div>
</div>
)}}
</div>
</div>
)
}

```

### 3.4 State Management

## Zustand Store Structure:

```
// store/vehicleStore.ts
import { create } from 'zustand'

interface VehicleState {
  vehicles: Vehicle[]
  selectedVehicle: Vehicle | null
  predictions: Prediction[]
  loading: boolean
  error: string | null

  // Actions
  fetchVehicles: () => Promise<void>
  selectVehicle: (vin: string) => void
  updateVehicleStatus: (vin: string, status: HealthStatus) => void
}

export const useVehicleStore = create<VehicleState>((set, get) => ({
  vehicles: [],
  selectedVehicle: null,
  predictions: [],
  loading: false,
  error: null,

  fetchVehicles: async () => {
    set({ loading: true })
    try {
      const response = await fetch('/api/vehicles')
      const vehicles = await response.json()
      set({ vehicles, loading: false })
    } catch (error) {
      set({ error: error.message, loading: false })
    }
  },

  selectVehicle: (vin) => {
    const vehicle = get().vehicles.find(v => v.vin === vin)
    set({ selectedVehicle: vehicle })
  },

  updateVehicleStatus: (vin, status) => {
    set(state => ({
      vehicles: state.vehicles.map(v =>
```

```
    v.vin === vin ? { ...v, health_status: status } : v
  )
  )))
}
```

## 3.5 Performance Optimization

### 3D Asset Optimization:

- Use Draco compression for GLTF models (70-90% size reduction)
- Implement LOD (Level of Detail) for complex models
- Use texture atlases to reduce draw calls
- Implement progressive loading (low-res → high-res)
- Use basis/KTX2 compressed textures

### Code Splitting:

- Route-based code splitting (Next.js automatic)
- Dynamic imports for heavy components
- Lazy load 3D models and charts
- Defer non-critical JavaScript

### Image Optimization:

- Next.js Image component with automatic WebP/AVIF
- Responsive images with srcset
- Lazy loading below-the-fold images
- CDN delivery via CloudFront

### Caching Strategy:

- Static pages cached at CDN edge
- API responses cached with SWR/React Query
- Service worker for offline support
- Optimistic UI updates

---

## 4. AGENTS & AI SYSTEMS

### 4.1 Agent Architecture

#### Agent Design Principles:

1. **Single Responsibility:** Each agent has one primary function
2. **Autonomy:** Agents make decisions within their domain without constant human intervention
3. **Observability:** All actions logged and monitorable
4. **Graceful Degradation:** Fallback to human operators when confidence is low
5. **Continuous Learning:** Agents improve from feedback loops

## 4.2 Master Agent Implementation

### Core Responsibilities:

- Central orchestrator coordinating all worker agents
- Maintains global state for all vehicles
- Enforces business rules and policies
- Routes tasks based on priority and agent availability
- Implements UEBA security checks
- Handles escalations

### Implementation Details:

```
# master_agent/orchestrator.py
class MasterAgent:
    def __init__(self):
        self.redis = Redis(host='redis-cluster', decode_responses=True)
        self.rabbitmq = RabbitMQConnection()
        self.ueba = UEBA SecurityLayer()

        # Worker agent registry
        self.workers = {
            'data_analysis': DataAnalysisAgent(),
            'diagnosis': DiagnosisAgent(),
            'customer_engagement': CustomerEngagementAgent(),
            'scheduling': SchedulingAgent(),
            'feedback': FeedbackAgent(),
            'quality_insights': QualityInsightsModule()
        }

    async def process_vehicle_event(self, vin: str, telemetry: dict):
        """Main event processing pipeline"""

        # Step 1: Log event
        event_id = await self.log_event(vin, telemetry)

        # Step 2: Security check
        if not await self.ueba.validate_data_source(telemetry):
```

```

        await self.raise_security_alert("Anomalous telemetry source", vin)
        return

    # Step 3: Update vehicle state
    await self.update_vehicle_state(vin, telemetry)

    # Step 4: Check if analysis needed
    if await self.should_analyze(vin):
        task = self.create_analysis_task(vin, telemetry)
        await self.rabbitmq.publish('data_analysis_queue', task)

    async def handle_prediction(self, prediction: dict):
        """Handle prediction from Diagnosis Agent"""

        vin = prediction['vin']
        confidence = prediction['confidence']
        priority = prediction['priority']

        # Business rules
        if confidence >= 0.85 and priority in ['P1', 'P2']:
            # High confidence, urgent: Initiate customer engagement
            await self.initiate_customer_engagement(prediction)

        elif confidence >= 0.70:
            # Medium confidence: Enhanced monitoring
            await self.schedule_enhanced_monitoring(vin)

        elif confidence >= 0.50:
            # Low confidence: Log for pattern analysis
            await self.log_low_confidence(prediction)

        else:
            # Very low confidence: Discard
            pass

    async def initiate_customer_engagement(self, prediction: dict):
        """Start customer engagement workflow"""

        vin = prediction['vin']

        # Get customer profile
        customer = await self.get_customer_profile(vin)

        # Check communication preferences

```

```

if not customer.consent_voice_call:
    # Fall back to app notification
    await self.send_app_notification(customer, prediction)
    return

# Check do-not-disturb hours
if await self.is_dnd_time(customer):
    # Schedule for later
    await self.schedule_delayed_engagement(customer, prediction)
    return

# Create engagement task
task = {
    'task_id': generate_uuid(),
    'vin': vin,
    'customer': customer.to_dict(),
    'prediction': prediction,
    'priority': prediction['priority'],
    'attempt': 1,
    'max_attempts': 3
}

# Publish to engagement queue
await self.rabbitmq.publish(
    'customer_engagement_queue',
    task,
    priority=self.get_priority_level(prediction['priority'])
)

# Update state
await self.update_vehicle_state(vin, {
    'engagement_status': 'INITIATED',
    'engagement_task_id': task['task_id']
})

```

### **State Machine Implementation:**

```

# master_agent/state_machine.py
from enum import Enum
from typing import Dict, Callable

class VehicleState(Enum):
    HEALTHY = "HEALTHY"
    MONITORING = "MONITORING"

```



```

PREDICTION_CONFIRMED = "PREDICTION_CONFIRMED"
CUSTOMER_CONTACTED = "CUSTOMER_CONTACTED"
APPOINTMENT_SCHEDULED = "APPOINTMENT_SCHEDULED"
SERVICE_IN_PROGRESS = "SERVICE_IN_PROGRESS"
SERVICE_COMPLETED = "SERVICE_COMPLETED"
FEEDBACK_COLLECTED = "FEEDBACK_COLLECTED"

```

```

class StateMachine:
    def __init__(self):
        self.transitions = {
            VehicleState.HEALTHY: {
                'anomaly_detected': VehicleState.MONITORING
            },
            VehicleState.MONITORING: {
                'prediction_confirmed': VehicleState.PREDICTION_CONFIRMED,
                'false_alarm': VehicleState.HEALTHY
            },
            VehicleState.PREDICTION_CONFIRMED: {
                'customer_contacted': VehicleState.CUSTOMER_CONTACTED,
                'contact_failed': VehicleState.PREDICTION_CONFIRMED # Retry
            },
            VehicleState.CUSTOMER_CONTACTED: {
                'appointment_booked': VehicleState.APPOINTMENT_SCHEDULED,
                'customer_declined': VehicleState.HEALTHY, # Log but move on
                'callback_requested': VehicleState.CUSTOMER_CONTACTED # Retry later
            },
            VehicleState.APPOINTMENT_SCHEDULED: {
                'service_started': VehicleState.SERVICE_IN_PROGRESS,
                'appointment_cancelled': VehicleState.PREDICTION_CONFIRMED # Reschedule
            },
            VehicleState.SERVICE_IN_PROGRESS: {
                'service_completed': VehicleState.SERVICE_COMPLETED
            },
            VehicleState.SERVICE_COMPLETED: {
                'feedback_received': VehicleState.FEEDBACK_COLLECTED
            },
            VehicleState.FEEDBACK_COLLECTED: {
                'cycle_reset': VehicleState.HEALTHY
            }
        }

```

```

async def transition(self, vin: str, event: str) -> bool:
    """Attempt state transition"""
    current_state = await self.get_current_state(vin)

```

```

if event not in self.transitions[current_state]:
    logger.warning(f"Invalid transition: {current_state} -> {event}")
    return False

new_state = self.transitions[current_state][event]

# Update state in Redis
await self.update_state(vin, new_state)

# Log transition
await self.log_transition(vin, current_state, new_state, event)

# Trigger side effects
await self.handle_state_change(vin, new_state)

return True

```

### 4.3 Data Analysis Agent

**Purpose:** Real-time anomaly detection from streaming telemetry

**Implementation:**

```

# agents/data_analysis_agent.py
import numpy as np
from sklearn.ensemble import IsolationForest
from kafka import KafkaConsumer
import asyncio

class DataAnalysisAgent:
    def __init__(self):
        self.consumer = KafkaConsumer(
            'vehicle.telemetry.raw',
            bootstrap_servers='kafka-cluster:9092',
            group_id='data_analysis_group',
            value_deser

```